

Questions & Answers on Chapter 01

? 1. List examples of real world applications of NLP

- Email platforms – e.g., Gmail spam detection
- Voice assistants – Siri, Alexa
- Search engines – Google understands queries
- Machine translation – Google Translate - Social media sentiment analysis
- Smart replies – Messaging apps
- Grammar checkers – Grammarly
- Chatbots – Customer support
- Healthcare – Extract info from clinical notes

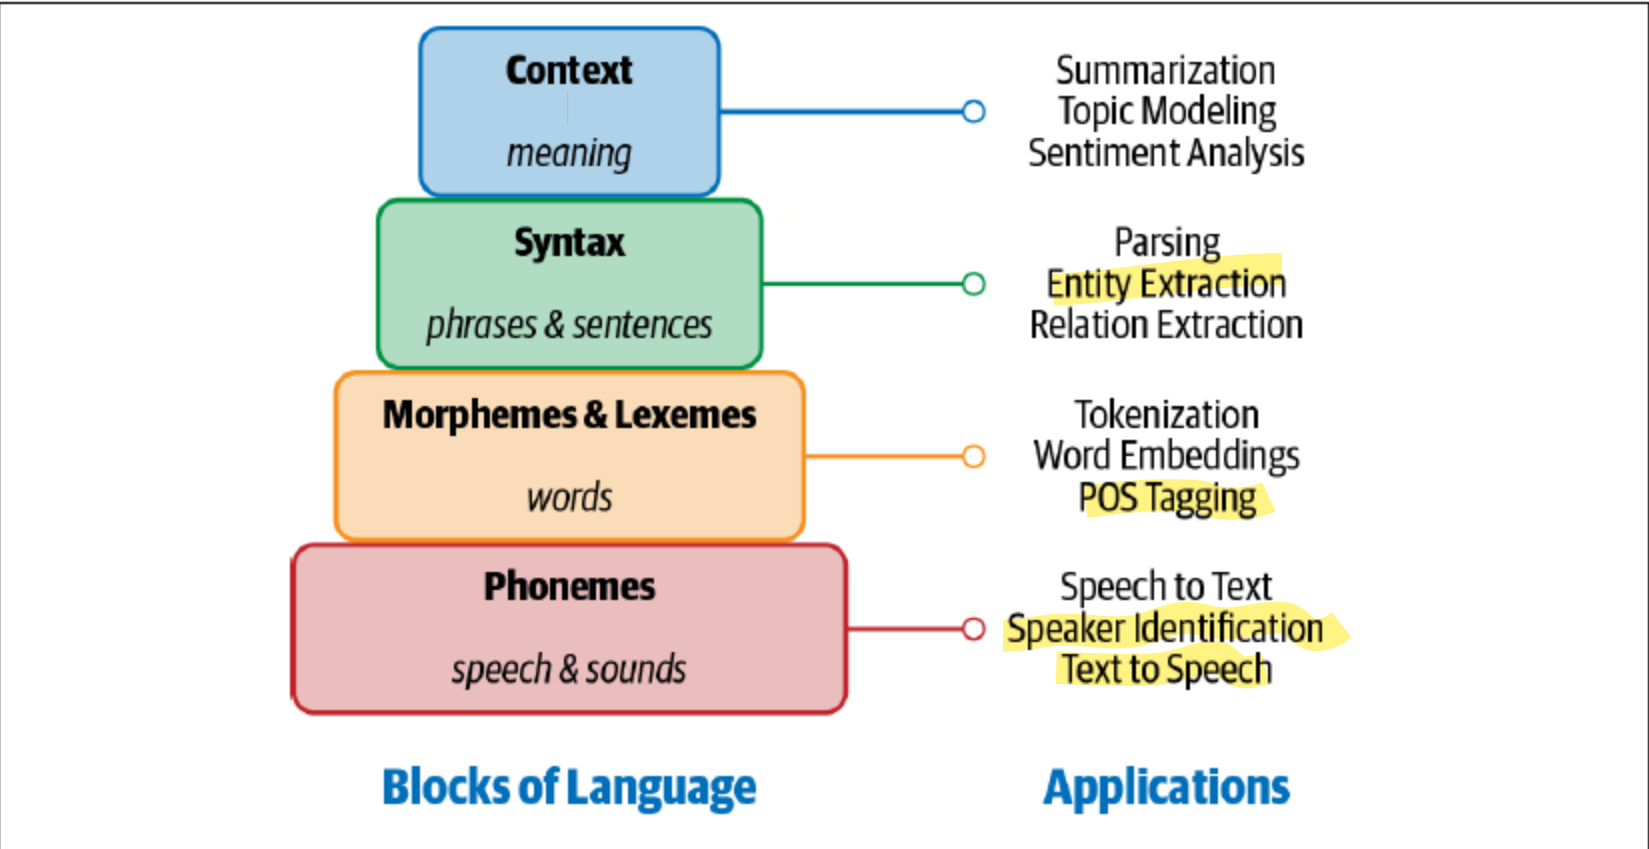
? 2. Explain the following NLP tasks: language modelling, text classification, information extraction, information retrieval, conversational agent, text summarization, question answering, machine translation, and topic modelling.

- Language Modelling: Predict the next word in a sequence.
- Text Classification: Label text (e.g., spam vs non-spam).
- Information Extraction: Pull entities (e.g., names, dates).
- Information Retrieval: Find relevant docs (e.g., Google).
- Conversational Agent: Systems that talk/understand (e.g., Siri).
- Text Summarization: Shorten long texts.
- ? Question Answering: Answer based on given input text. Watson AI
- Machine Translation: Translate across languages.
- Topic Modelling: Find topics/themes in large text collections.
Group docs by themes

? 3. What are the building blocks of language and their applications?

Smallest sound units (e.g., "sh" in "shoe").

- Phonemes – Smallest sound units → Used in speech-to-text
- Morphemes / Lexemes – Meaningful units → Tokenization, embeddings
- Syntax – Sentence structure → Parsing & grammar checking
- Context – Real-world meaning → NER, summarization



? 4. Why is NLP Challenging?

- 🗣️ **Ambiguity** – Words have multiple meanings
- 🧠 **Lack of common sense** – Machines lack human knowledge
- 🎨 **Creativity** – Language is playful, poetic, diverse → Idioms, slang, poetry confuse models.
- 🌐 **Multilingual complexity** – Different rules, scripts, dialects
 ↳ Diverse languages: Rules vary across languages.

? 5. How NLP, ML, and DL are related?

- 🧠 **AI** – The parent field
- 🤖 **ML** – Subfield: learns from data
- ⚡ **DL** – Subfield of ML: deep neural nets
- 🗣️ **NLP** – Uses ML/DL for understanding human language

? 6. Describe the heuristics-based NLP.

- 🛠️ Based on rules or word counts
 Count "positive" words to guess sentiment.
- Example: Sentence with “good” = positive 😊
- ✅ Simple
- ❌ Not flexible, hard to scale

? 7. Explain briefly Naive Bayes, Support Vector Machine, Hidden Markov Model, And Conditional Random Fields approaches.

- 🧪 **Naive Bayes**: fast but assumes features are independent. Probabilistic classification. Uses probability & Bayes’ rule for classification.
- ⚙️ **SVM**: Draws a boundary between classes; great for high-dimensional data. robust but slow
- 🔒 **HMM**: Probabilistic model for sequences; Models hidden states great for speech tagging.
- 📁 **CRF**: Like HMM but more powerful; good for sequence labeling (NER, POS).
 Like HMM but better! Tags each word in context (e.g., "Apple" = company or fruit?).

? 8. What is the difference between RNN and LSTM NN?

- 🔄 **RNN**: Processes sequences; Remembers short-term steps.
- ⌚ **LSTM**: Advanced RNN that remembers long-term context better (uses gates 🧠 🚪)
 Remembers important info, forgets irrelevant details.

? 9. How CNN can be used for text processing?

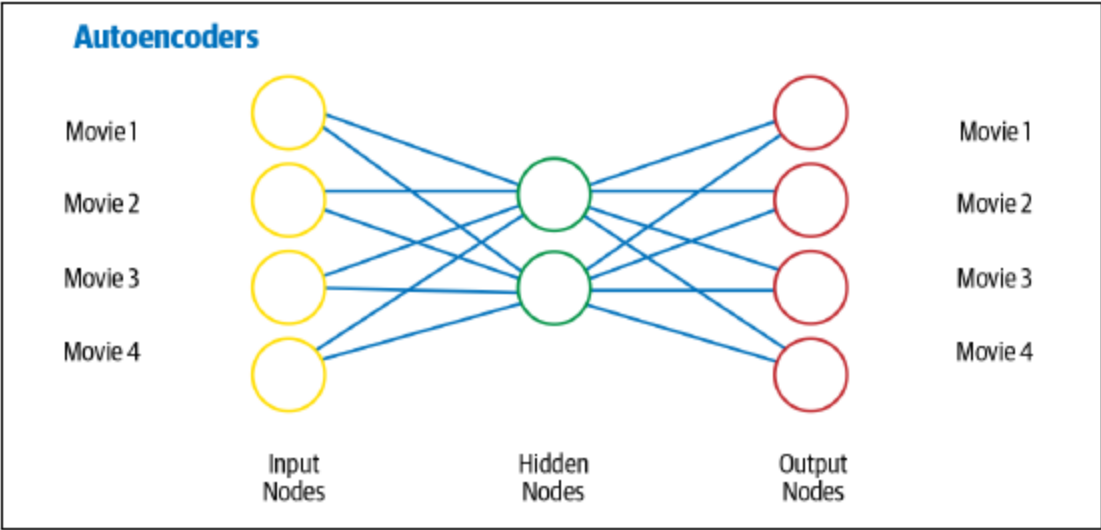
- 🏗️ Convert text to vectors
- 🧠 CNN identifies patterns (e.g., n-grams)
- ✅ Fast, great for classification tasks

? 10. Describe the concept transfer learning. Remember Transformers

- 📚 Pre-train on big data (e.g., BERT), then fine-tune for specific tasks.
- 🖼️ Train on one task → fine-tune for another
- Example: Pretrain on Wikipedia → Fine-tune for medical docs 📄
- ✅ Works well when labeled data is small

? 11. Give the architecture of autoencoder.

- ➡️ **Encoder** – Compress input
- 🔗 **Bottleneck** – Latent vector
 Input -> Encoder (compresses) -> Latent vector -> Decoder (reconstructs).
- ⬅️ **Decoder** – Reconstructs original
- 🎯 Application: Feature learning, denoising

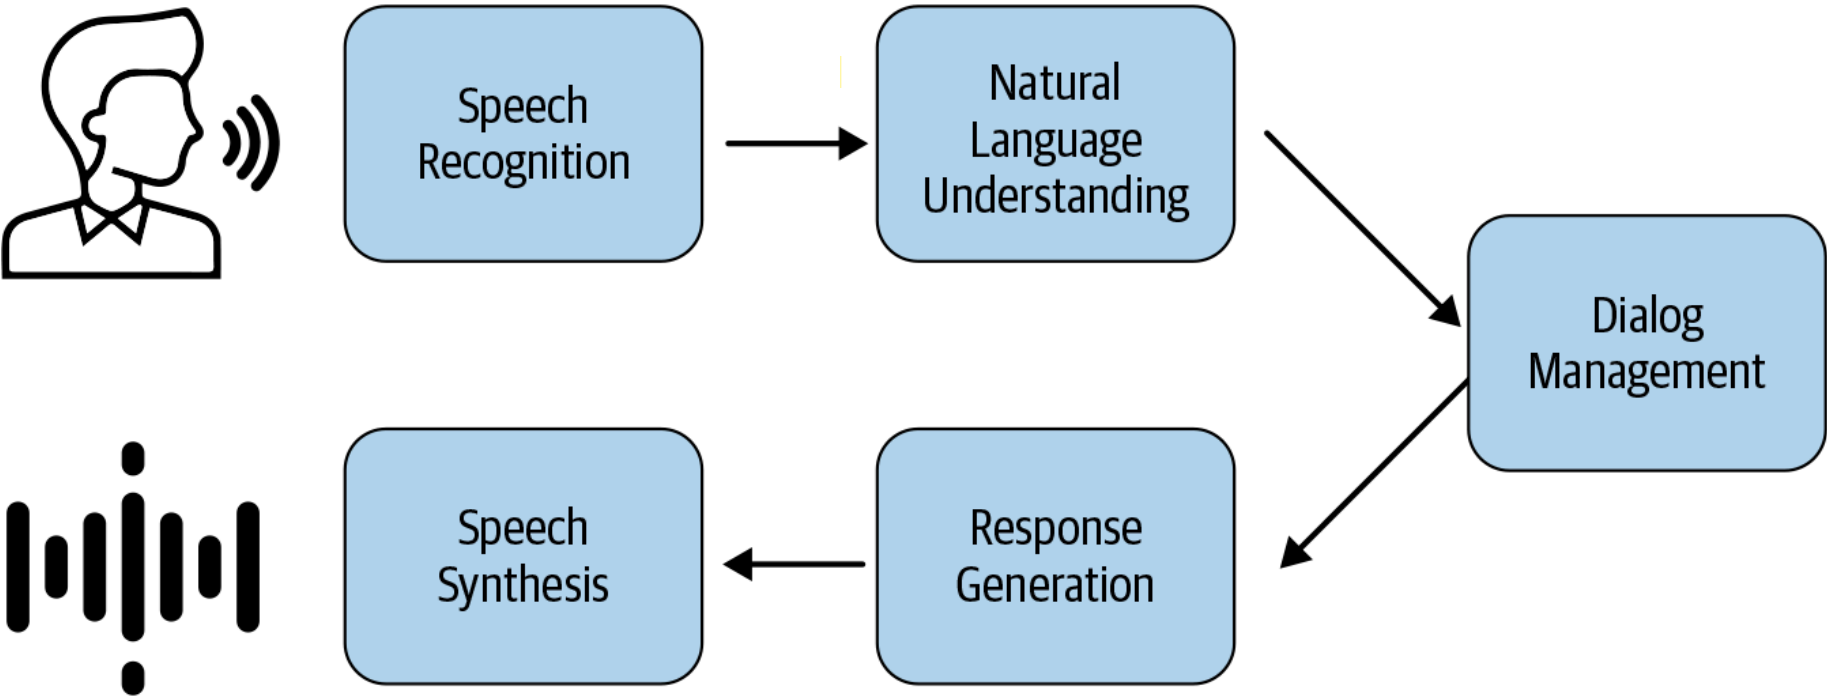


? **12. List the key reason that makes DL not suitable for all NLP tasks**

- Needs lots of data → Overfitting on small sets
- Expensive to train/deploy
- Hard to interpret
- Poor transferability to new domains Fails with domain shifts
- Not efficient on low-power devices





? **13. Explain the flow of conversation agents**

- Speech Recognition: Convert voice to text
- Language Understanding
 - Sentiment analysis
 - Named Entity Recognition (NER)
 - Coreference resolution (who is “he”?)
- Dialog Management – Understand intent, decide next step
- Response Generation – Generate a meaningful reply
- Text-to-Speech – Speak the response (optional)





Questions & Answers on Chapter 03

? 1. List the four categories of text representation techniques.

-  **Basic vectorization**: One-hot, BoW, TF-IDF
-  **Distributed representations**: Word2Vec, GloVe
-  **Universal language representations**: BERT, ELMo
-  **Handcrafted features**: Custom, domain-based

? 2. Describe the concept vector space models.

-  A **Vector Space Model (VSM)** represents text as numerical vectors.
-  These are inputs to ML models (e.g., for classification or similarity).

✦ *Example:*

“Dog bites man” → [1, 1, 1, 0, 0, 0]

? 3. Use “D1: Dog bites man, D2: Man bites dog, D3: Dog eats meat, and D4: Man eats food” as an input, find their representation using one-hot encoding, bag of words, bag of N-gram, and TF-IDF.

◆ Vocabulary

First, build the vocabulary from all four docs:

1 2 3 4 5 6

Vocab: [dog, bites, man, eats, meat, food]

Size of vocabulary (|V|) = 6

TF (term-freq) = no.of term occurred in a doc / total no.of terms
IDF = log(total no.of docs / no.of docs contain the term)
TF X IDF

✦ 1. One-Hot Encoding

Each word = a unique binary vector of size |V| = 6

Word	One-hot Vector
dog	[1, 0, 0, 0, 0, 0]
bites	[0, 1, 0, 0, 0, 0]
man	[0, 0, 1, 0, 0, 0]
eats	[0, 0, 0, 1, 0, 0]
meat	[0, 0, 0, 0, 1, 0]
food	[0, 0, 0, 0, 0, 1]

Document representations (each word has its own row):

- D1 → dog = [1, 0, 0, 0, 0, 0] , bites = [0, 1, 0, 0, 0, 0] , man = [0, 0, 1, 0, 0, 0]

➡ Simple but sparse, no semantic meaning, and can't handle unknown words (OOV).

🛒 2. Bag of Words (BoW)

We count how many times each word appears in a document. Order doesn't matter ✖ 

Doc	dog	bites	man	eats	meat	food
D1	1	1	1	0	0	0
D2	1	1	1	0	0	0
D3	1	0	0	1	1	0
D4	0	0	1	1	0	1

- Pros: Easy, fixed-length, works well with simple models
- Cons: No order, no context, high-dimensional

🔗 3. Bag of N-Grams (e.g., Bigrams = 2-grams)

Make sequences of 2 consecutive words:

- D1: dog bites , bites man
- D2: man bites , bites dog
- D3: dog eats , eats meat
- D4: man eats , eats food

Vocabulary of Bigrams: [dog bites, bites man, man bites, bites dog, dog eats, eats meat, man eats, eats food]

Doc	dog bites	bites man	man bites	bites dog	dog eats	eats meat	man eats	eats food
D1	1	1	0	0	0	0	0	0
D2	0	0	1	1	0	0	0	0
D3	0	0	0	0	1	1	0	0
D4	0	0	0	0	0	0	1	1

- Captures some word order and context, better than BoW
 - Still sparse and doesn't handle meaning
- TF (term-freq) = no.of term occurred in a doc / total no.of terms in doc
IDF = $\log_2(\text{total no.of docs} / \text{no.of docs contain the term}) = \log_2(N/DF)$
DF -> Document Freq
TF X IDF

📊 4. TF-IDF (Term Frequency – Inverse Document Frequency)

TF = How often a word appears in a doc
IDF = How rare the word is across all docs

For example:

- “dog” appears in D1, D2, D3 → common → lower weight
- “meat” and “food” appear only once → higher weight

TF (Term Frequency)

Word	D1	D2	D3	D4
dog	0.33	0.33	0.33	0
bites	0.33	0.33	0	0
man	0.33	0.33	0	0.33
eats	0	0	0.33	0.33
meat	0	0	0.33	0
food	0	0	0	0.33

Resulting TF-IDF vector (simplified):

Word	TF-IDF (D1)
dog	low
bites	medium
man	low
eats	0
meat	0
food	0

IDF (Inverted Doc Frequency)

Word	DF	IDF
dog	3	0.41
bites	2	1
man	3	0.41
eats	2	1
meat	1	2
food	1	2

TF-IDF (TF(for word in each doc) X IDF)

Word	D1	D2	D3	D4
dog	0.33*0.41	0.136	0.136	0
bites	0.33	0.33	0	0
man	0.136	0.136	0	0.136
eats	0	0	0.33	0.33
meat	0	0	0.66	0
food	0	0	0	0.66

- TF-IDF gives importance to rare words and downweights common ones
- ✓ Great for document similarity and ranking

? 4. Explain the difference between:

- (a) distributional similarity and distributional hypothesis
- (b) distributional representation and distributed representation

- ♦ (a)

اختلافی

 - **Distributional hypothesis**: Words in similar contexts have similar meanings.
 - **Distributional similarity**: Measures context-based word similarity.
- ♦ (b)

مقبول

ex: BOW, co-occurrence matrix

 - **Distributional representation**: High-dim vectors from co-occurrence.
 - **Distributed representation**: Dense, learned vectors (e.g., Word2Vec, GloVe).

low-dimensional vectors (small)

? 5. Describe the **wording embedding** concept with an example of its use.

- 📦 **Word embeddings**: Dense **vectors encoding meaning**.
 - ➡ **Similar words = close in vector space**.
- 🔴 *Example:*
"USA" is close to "Canada", "Germany"
- EX: In sentiment analysis, embeddings help models see that "great" = "excellent" = "fantastic", so different words with similar meaning give similar results.

? 6. Explain with an example the **two architectural variants of Word2vec: CBOW and SkipGram**.

- ✖ **CBOW**: **Predict center word** from context
- 🧠 **SkipGram**: **Predict context** from center word

- 🔴 *Example:*
Sentence: "The cat sat on the mat"
- CBOW: ("The", "sat") → "cat"
 - SkipGram: "cat" → ("The", "sat")

Out-of-Vocabulary

? 7. How the **OOV problem** can be solved?

- ⚙ **Subword models** (prefix/suffix)
- abc **Character-level** or **byte-level** models
- ⚙ **Random init** for unknown words
- 💬 Use **fastText** (**subword-based** vectors)

? 8. What is the difference between **Doc2vec** and **Word2vec**?

- 🧠 **Word2Vec**: **Word-level** vectors
- 📄 **Doc2Vec**: **Document-level** vectors
- ➡ Doc2Vec = **Word2Vec + document context**

? 9. What are the **important aspects to keep in mind while using word embeddings**?

- ⚠ Bias can leak from data
- 📦 Embeddings = **large file sizes**
- 💡 **Not all linguistic structure is captured** ←
- 🧠 **Might need extra sentence/document-level features**

? 10. **How high-dimensional data can be represented visually?**

- 📊 Use **t-SNE**:
 - **Reduces high-dim embeddings** (e.g., 300D) → 2D/3D
 - **Useful for visualizing clusters/semantic similarity**

المستطابق في التوزيع

? 11. With example explain the **use of handcrafted** feature representations.

- 🧠 Example: **TextEvaluator** (by ETS)
- 📖 Used for **estimating reading difficulty**
- **Handcrafted features include:**

- Syntactic complexity
- Word familiarity
- Concreteness

منعقدة
و حسي

Useful for educational and grading tasks.

إِنَّ اللَّهَ وَمَلَائِكَتَهُ يُصَلُّونَ عَلَى النَّبِيِّ يَا أَيُّهَا الَّذِينَ آمَنُوا صَلُّوا عَلَيْهِ وَسَلِّمُوا تَسْلِيمًا (56)

Questions & Answers on Chapter 04

? 1. What is the difference between binary, multi-class, and multi-label classification?

- **Binary:** Two classes only (e.g., spam vs. not spam)
- **Multiclass:** More than two classes, choose one (e.g., positive / neutral / negative)
- **Multilabel:** One text can have multiple labels (e.g., news tagged as “sports” and “politics”)

? 2. Give some applications of text classification.

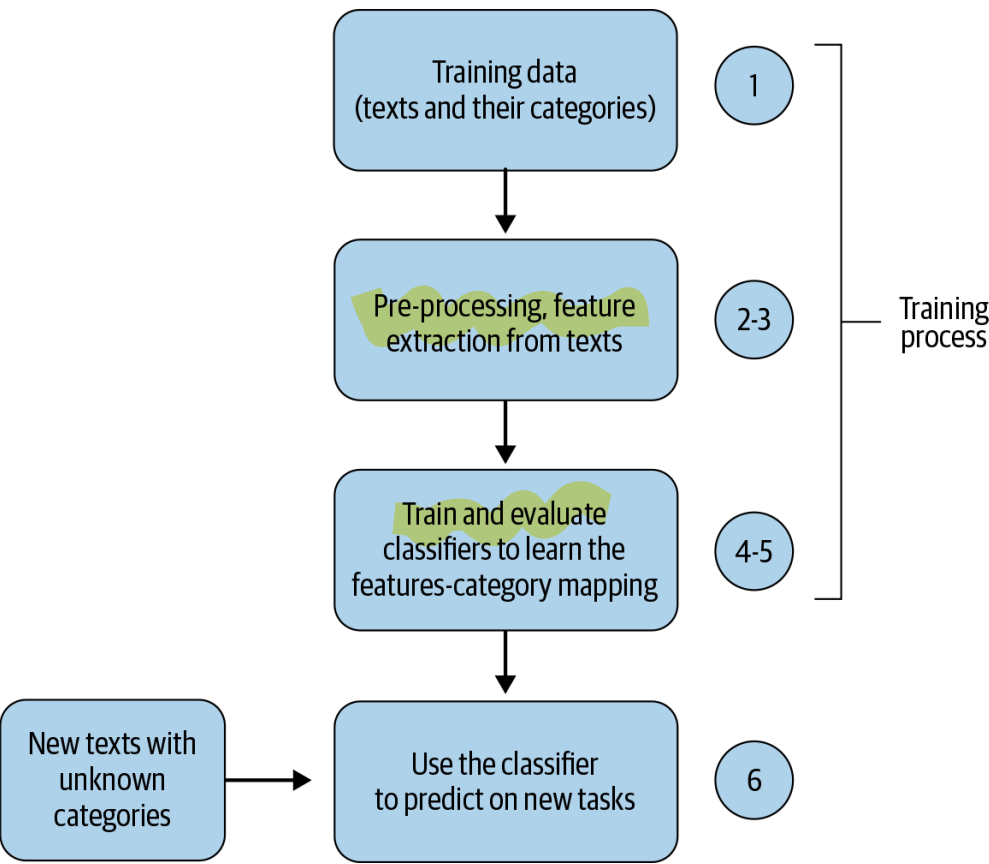
Classifying articles, blogs, or news into topics (e.g., sports, tech)

- 📁 Organizing content (e.g., news, blogs)
- 🗣️ Customer support automation
- 🛒 Sentiment analysis in e-commerce Detecting positive/negative reviews about products
- 🗣️ Language identification Detecting the language of the text (English, Arabic, French...)
- 📰 Fake news detection Identifying whether a news article is real or fake

- Spam filtering -> Classifying emails as "Spam" or "Not Spam"

? 3. Describe the pipeline for building text classification systems.

1. Collect labeled data
2. Split into train/test/validation sets
3. Convert text to features (vectors)
4. Train the model 🧠
5. Evaluate using metrics (accuracy, F1, etc.)
6. Deploy the model 🚀 and monitor performance



? 4. Classification can be done without the text classification pipeline, explain how?

dictionary

You can use a **lexicon-based approach**:

- Make lists of positive and negative words
- Use rules (e.g., 😊 = positive, 😞 = negative)

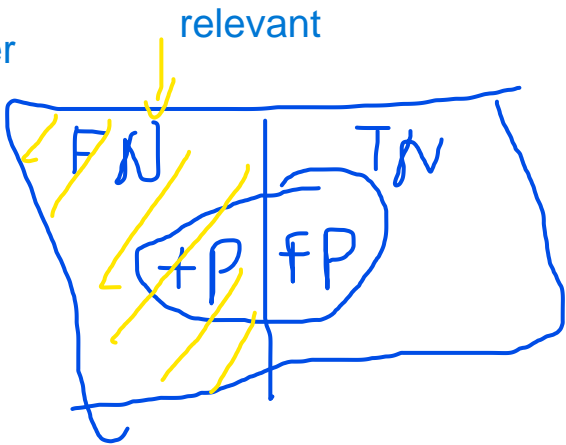
Simple, but limited and not scalable for complex tasks.

? 5. Describe with an example the confusion matrix of a classifier.

A confusion matrix helps measure performance by comparing actual vs predicted labels.

Example (binary classification): ex: Spam Filter

	Predicted 📧	Predicted 📧
Actual 📧	True Pos (TP)	False Neg (FN)
Actual 📧	False Pos (FP)	True Neg (TN)



Metrics:

- Precision = TP / (TP + FP) tp / retrieved (tp+fp)
- Recall = TP / (TP + FN) tp / relevant (tp+fn)

? 6. List the potential reasons for poor classifier performance.

- Wrong or weak algorithm The chosen model is too simple or not suitable
- Poor pre-processing / features Text not cleaned properly, or bad feature extraction
- Bad hyperparameter tuning Parameters like learning rate, tree depth, etc., are not optimized
- Imbalanced data or noisy labels One class has way more data than others
 - Noisy or incorrect labels Training data contains wrong or inconsistent labels
 - Too little training data

? 7. How to solve class imbalance problem of a dataset?

- Use balanced metrics (e.g., F1 score) Instead of accuracy, use F1-score, precision, recall
- Resample data (oversample or undersample) SMOTE
- Try class-weighted models give more weight to the minority class
- Train custom models that adjust for imbalance

? 8. What is the difference between generative and discriminative classifiers?

- Generative: Models how data is generated → calculates P(x|y) and P(y) Generator => models the world (e.g., Naive Bayes)
- Discriminative: Directly models decision boundary → learns P(y|x) Decision => draws the line between classes (e.g., Logistic Regression, SVM)

Dense vectors encoding meaning. Similar words = closer in vector space

? 9. How to use word embeddings as features for text classification?





- 1 Replace words with pre-trained vectors (Word2Vec, GloVe) each word replaced with vector
- 2 Average them or feed them into a neural net convert to numeric representation
- ✓ Captures meaning and similarity → using model
- ✦ Used for tasks like sentiment classification

? 10. List the steps for converting training and test data into a format suitable for the neural network.

Converts text into machine-readable words/tokens.


1. Tokenize the sentences
2. Pad them to same length [1, 2, 3] -> [1, 2, 3, 0] (with padding)
3. Map words to embedding vectors Represent words as meaningful dense vectors.
4. Feed embeddings into neural net (CNN, LSTM, etc.) Learns patterns and performs classification.


? 11. Which technique is better for text classification CNN or LSTM and why?

-  Use **LSTM** for long sequences or large datasets (remembers order)
-  Use **CNN** for shorter text and faster training
-  LSTM = better for capturing long-term dependencies
-  CNN = good at local patterns (like n-grams)

Use LSTM when word order matters or text is long.
Use CNN when you need speed and are okay with losing some sequence info.



? **12. How text classification models can be interpreted?**

Use **LIME**  (Local Interpretable Model-Agnostic Explanation):






- Explains predictions word by word
- Shows which features affected the decision
-  Useful in sensitive tasks (e.g., medical, legal)
 - Medical predictions
 - Legal document classification
 - Any critical decision-making task

Example:
If a classifier predicts a review is positive, LIME might highlight:
"I love this phone" => LIME shows "love" had the biggest impact

? **13. How to solve no training and less training data problems?**

-  **No data** → Create dataset by annotation manual annotation (label the data yourself or with help)
-  **Less data** → Use:
 - **Active learning** Let the model choose the most useful examples to label
 - **Domain adaptation** Use models trained on similar tasks/domains
 - **Weak supervision** (rules or heuristics) Use rules, patterns, or heuristics to auto-label data

? **14. Give some options to explore when no labels exist for a dataset.**

-  Use pre-built APIs (Google NLP, etc.)
-  Find public labeled datasets
-  Apply weak supervision
-  Try active learning
-  Learn from feedback over time

? **15. Describe the pipeline for building a classifier when there is no training data.**

1. Define rules/heuristics "Love" -> pos, "bad" -> neg , ...
2. Use **Snorkel** or other tools for weak labeling
3. Leverage existing models/APIs (Google NLP – HuggingFace models)
4. Apply active learning to label selectively
5. Iterate and refine with real user feedback

(56) إِنَّ اللَّهَ وَمَلَائِكَتَهُ يُصَلُّونَ عَلَى النَّبِيِّ يَا أَيُّهَا الَّذِينَ آمَنُوا صَلُّوا عَلَيْهِ وَسَلِّمُوا تَسْلِيمًا

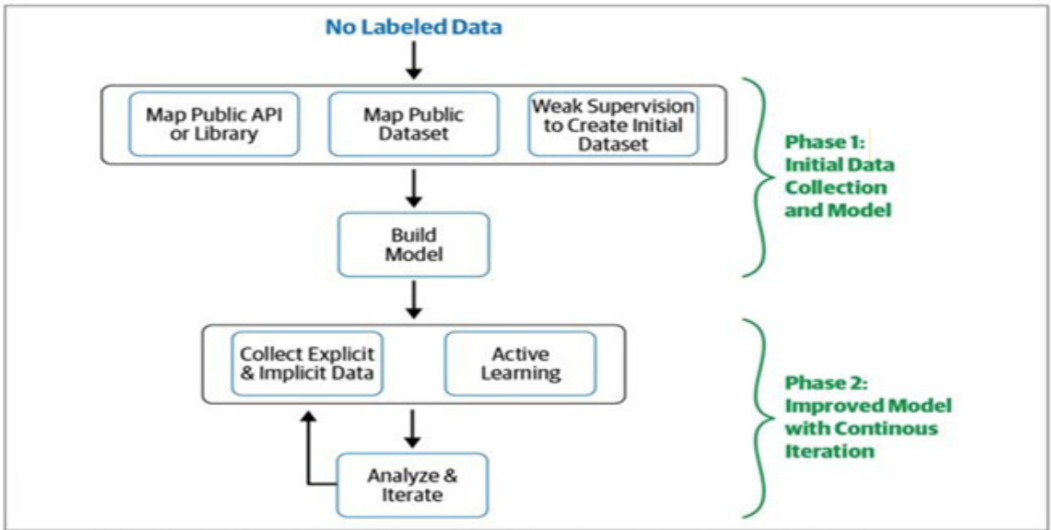
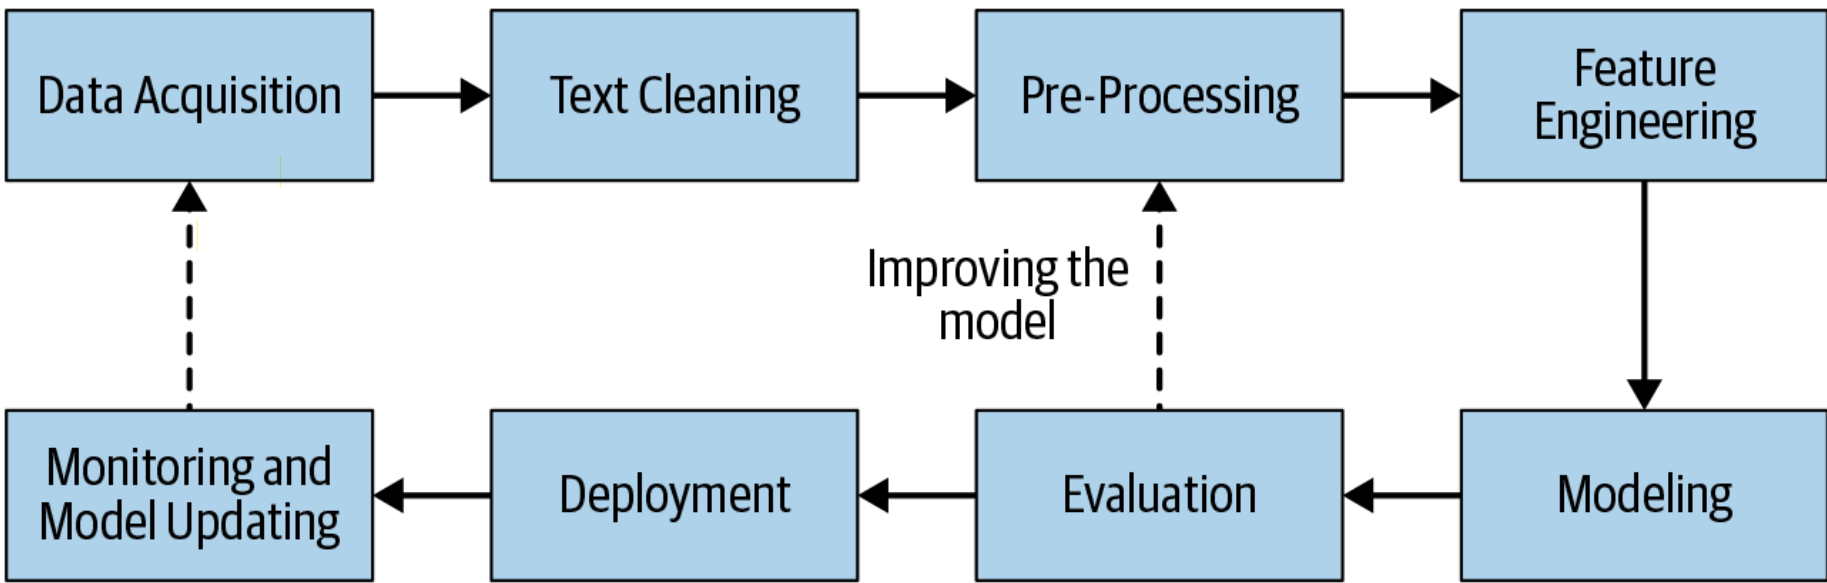


Figure 4-11. A pipeline for building a classifier when there's no training data

Questions & Answers on Chapter 02

? 1. What are the key stages of a generic pipeline for NLP system development?

- 🛒 **Data acquisition** → Get the data
- 🧼 **Text cleaning** → Remove noise like HTML, symbols
- 📄 **Pre-processing** → Tokenize, lowercase, etc.
- 🧬 **Feature engineering** → Convert text into numbers
- 🧠 **Modeling** → Train the ML/DL model
- 🎯 **Evaluation** → Test performance with metrics
- 🚀 **Deployment** → Integrate into a product
- 📊 **Monitoring & Updating** → Improve with feedback



? 2. How can we get data required for training an NLP technique?





- ✅ Use public datasets (Kaggle, Hugging Face)
- 🕷 **Scrape websites** *intervention*
- 📱 **Product instrumentation** (collect from user input)
- 📄 **Data augmentation:**
 - 🔄 **Synonym replacement** ("happy" => "joyful")
 - 🌐 **Back translation** (Translate EN => FR => EN)
 - 🔄 **Bigram flipping** ("machine learning" => "learning machine")
 - 🧑 **Entity replacement**
 - 🗣 **Add noise** ("hello" => "hlelo")

? 3. List the different data augmentation methods

- 🔄 **Synonym replacement** – Swap words with similar ones
- 🌐 **Back translation** – Translate text to another language then back
- ↔ **Bigram flipping** – Swap pairs of words
- 🧑 **Entity replacement** – Change names/places with others
- 🗣 **Add noise** – Introduce typos or symbols randomly
- 🧰 **Tools:**
 - 🔄 **Snorkel** – Weak supervision
 - ⚙ **EDA, NLPAug** – Generate synthetic data
 - 🎯 **Active learning** – Label the most useful samples

? 4. Data can be collected from PDF files, HTML pages, and images. How can this data be cleaned based on their sources?

Use PyPDF2 to extract text.

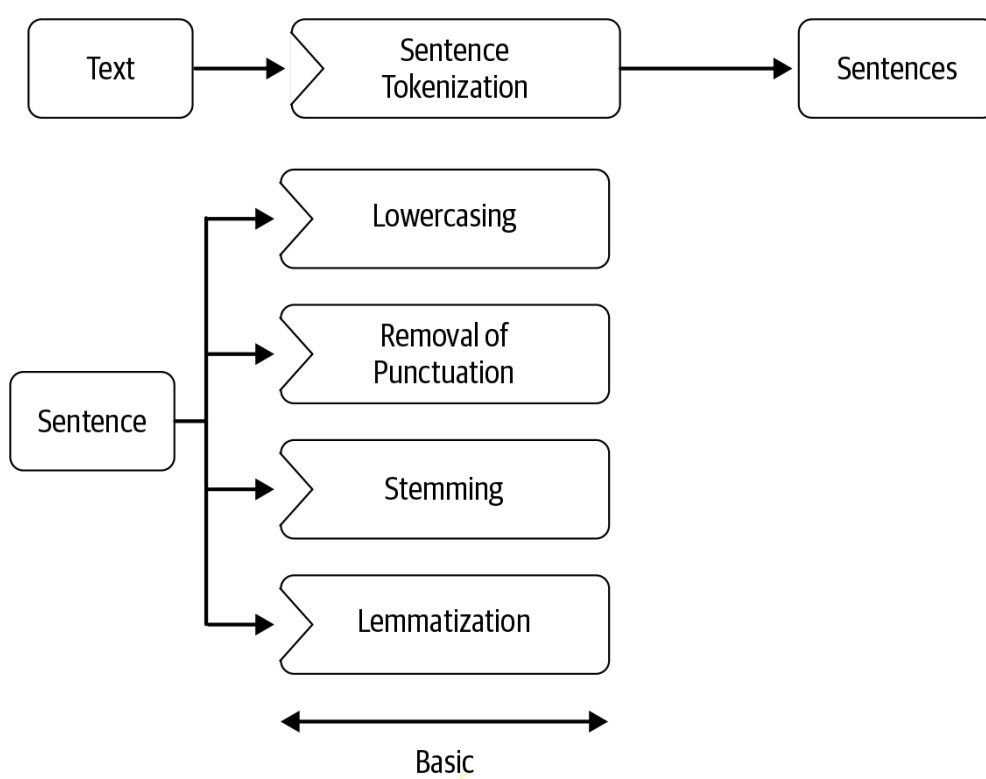
-  **PDFs** → Use PyPDF/Text extraction
-  **HTML** → Strip tags using **BeautifulSoup**
-  **Images** → Use **OCR** (e.g., Tesseract)
-  **Then:** remove formatting, special characters, fix encodings

? 5. Using dot (.) to segment sentences can cause problems, explain how?

- **Problem:** Splits abbreviations (e.g., "Mr. Smith" → "Mr" + ".").
- **Solution:** Use NLP libraries (spaCy, NLTK) to detect context.


? 6. What are the frequent steps in the data pre-processing phase?

1. Stop word removal 🚫 ("the", "is").
2. Stemming/Lemmatization 🌱 ("running" → "run").
3. Remove digits/punctuation 🗑️ ("Price: \$100" → "price").
4. Lowercasing 🔡 ("Hello" → "hello").



? 7. With examples, explain the differences between segmentation and lemmatization

Splitting text into sentences/words

 **Segmentation** – Break text into parts

→ “Hi. I’m Anas.” → 2 sentences

Lemmatization – Reduce to base form

→ “Was” → “be”, “Better” → “good”

? 8. What is the difference between code mixing and transliteration?

Mixing languages in one sentence

 Code mixing – Switching languages mid-sentence

→ “*Ana hungry w 3ayez akol*”

Transliteration – Non-English written using English letters

→ “shukran” instead of شكرًا

9. Describe the concept of coreference resolution

Link pronouns to their nouns.

 Coreference resolution links words to the same entity

→ “Sara went home. She was tired.” → “She” = “Sara”

📌 Used in summarization, QA, info extraction

? 10. Explain the feature engineering for classical NLP vs DL-based NLP

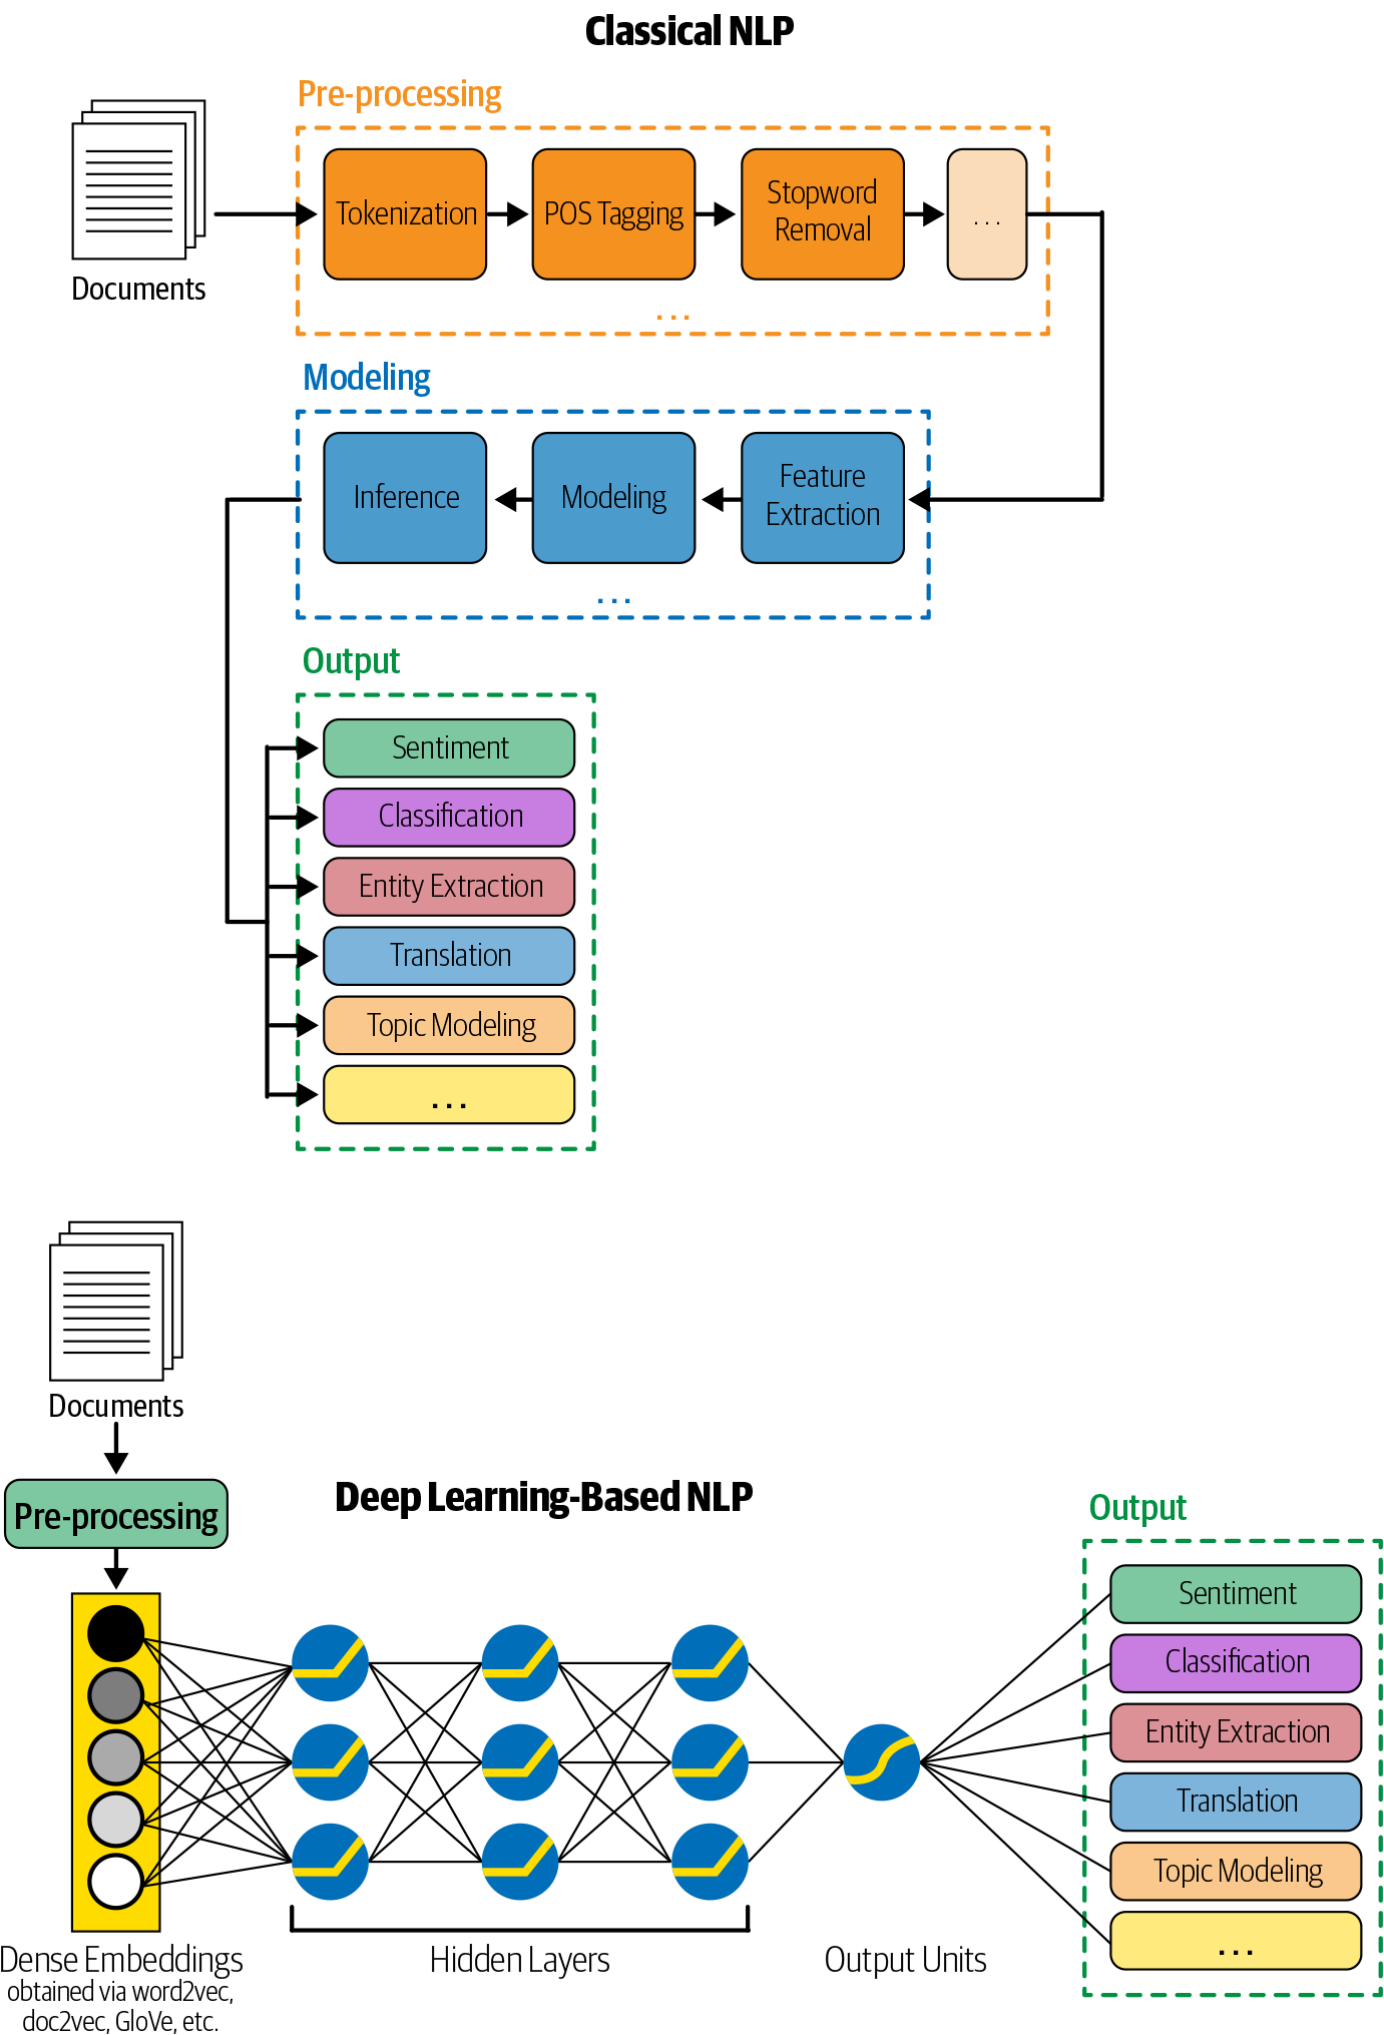
🔧 **Classical NLP:** Example: Counting "positive" words for sentiment analysis

Handcrafted features

- Manual features (BoW, TF-IDF, sentiment, etc.)

⚡ **Deep Learning:** (BERT embeddings)

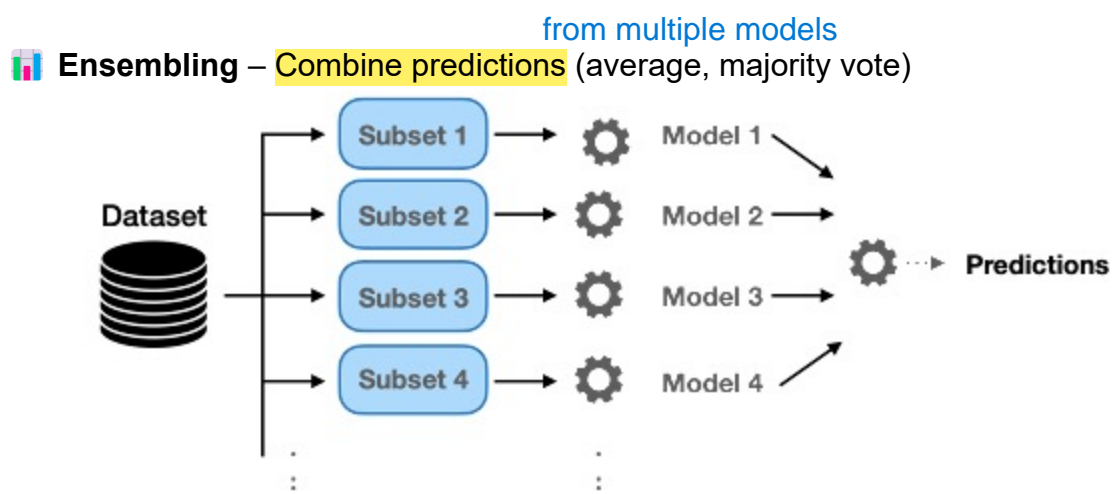
- Raw text → embeddings → automatic feature learning




? 11. How to combine heuristics directly or indirectly with the ML model?

- ♦ **As features** – Add them to model input
- ♦ **Preprocessing** – Use to filter or convert data
- ✓ Best when heuristics are reliable!

? 12. What is the difference between model ensembling and stacking?







 **Stacking** – Use one model's output as input to another



Example: Feed SVM predictions into a logistic regression model










? **13. Which modeling technique can be used in the following cases:** small data, large data, poor data quality, and good data quality ?

-  **Small data** → Rule-based or traditional ML
-  **Large data** → Deep learning
-  **Poor data** → Clean more first
-  **Good data** → Use off-the-shelf models quickly
Use cloud APIs (e.g., Google NLP).




? **14. What is the difference between intrinsic and extrinsic evaluation?**

-  **Intrinsic** – Direct metrics (e.g., accuracy, F1) numbers
-  **Extrinsic** – Impact on task (e.g., user satisfaction), resolution time




? **15. What are the metrics that can be used in:** classification, measuring model quality, information retrieval, predication, machine translation, and summarization tasks ? 

-  **Classification** → Accuracy, F1- score
-  **Model quality** → AUC Confusion Matrix
-  **Info retrieval** → MRR, MAP
-  **Prediction** → RMSE (for regression)
-  **Machine translation** → BLEU, METEOR
-  **Summarization** → ROUGE

? **16. Describe** deploying, monitoring, and updating phases of NLP pipeline

-  **Deployment** – Expose model as service/API
-  **Monitoring** – Track feedback, quality drift Track accuracy drops
-  **Updating** – Periodically retrain on new data

? **17. Explain how the NLP pipeline is different from one language to another**

-  **English** – Easier segmentation, fewer morphology issues
-  **Chinese/Arabic** – Need custom tokenizers, more complex grammar
-  Different languages may require different tools, steps, and models

? 18. Describe the NLP pipeline for ranking tickets in a ticketing system by Uber

- 📄 Collect ticket → Clean HTML
- ✂️ Preprocess – tokenize, lowercase, remove stop words, lemmatize
- 🖼️ Feature engineering:
 - BoW, TF-IDF, LSI Topic Modeling
 - Compare ticket & topics with cosine similarity
- 🧠 Modeling – Use ticket + trip data
- 🎯 Ranking – Select top 3 solutions
- 📊 Evaluation:
 - 📊 Intrinsic – MRR
 - 🎯 Extrinsic – Customer satisfaction, resolution time

